


IF3140 – Sistem Basis Data Transactions

SEMESTER I TAHUN AJARAN 2024/2025



KNOWLEDGE & SOFTWARE ENGINEERING



Summer

Silberschatz, Korth, Sudarshan:
“Database System Concepts”,
7th Edition

- Chapter 17: Transactions

Objectives

Students are able to:

- Explain the importance of transaction properties
- Explain serializable transactions
- Explain the concept of implicit commits
- Describe the issues specific to efficient transaction execution
- Explain at least two transaction protocols



Outline

- Transaction Concept
- Transaction State
- Concurrent Executions
- Serializability
- Recoverability
- Implementation of Isolation
- Transaction Definition in SQL
- Testing for Serializability



Transaction Concept

A **transaction** is a *unit* of program execution that accesses and possibly updates various data items.

Example

```
1.read(A)
2.A := A - 50
3.write(A)
4.read(B)
5.B := B + 50
6.write(B)
```

Main issues

- Failures of various kinds
- Concurrent execution of multiple transactions

Example of Fund Transfer

Transaction to transfer \$50 from account A to account B:

1. **read**(A)
2. $A := A - 50$
3. **write**(A)
4. **read**(B)
5. $B := B + 50$
6. **write**(B)

- **Atomicity requirement**

- if the transaction fails after step 3 and before step 6, money will be “lost” leading to an inconsistent database state
 - Failure could be due to software or hardware
- the system should ensure that updates of a partially executed transaction are not reflected in the database

- **Durability requirement** — once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place), the updates to the database by the transaction must persist even if there are software or hardware failures.



Example of Fund Transfer (Cont.)

Transaction to transfer \$50
from account A to account B:

1. **read**(A)
2. $A := A - 50$
3. **write**(A)
4. **read**(B)
5. $B := B + 50$
6. **write**(B)

- **Consistency requirement** in example:
 - The sum of A and B is unchanged by the execution of the transaction
 - In general, consistency requirements include:
 - Explicitly specified integrity constraints such as primary keys and foreign keys
 - Implicit integrity constraints
 - e.g. sum of balances of all accounts, minus sum of loan amounts must equal value of cash-in-hand
 - A transaction must see a consistent database.
 - During transaction execution the database may be temporarily inconsistent.
 - When the transaction completes successfully the database must be consistent
 - Erroneous transaction logic can lead to inconsistency



Example of Fund Transfer (Cont.)

T_1	T_2
1. read(A)	
2. $A := A - 50$	
3. write(A)	read(A), read(B), print(A+B)
4. read(B)	
5. $B := B + 50$	
6. write(B)	

Isolation requirement — if between steps 3 and 6, another transaction T_2 is allowed to access the partially updated database, it will see an inconsistent database (the sum $A + B$ will be less than it should be).

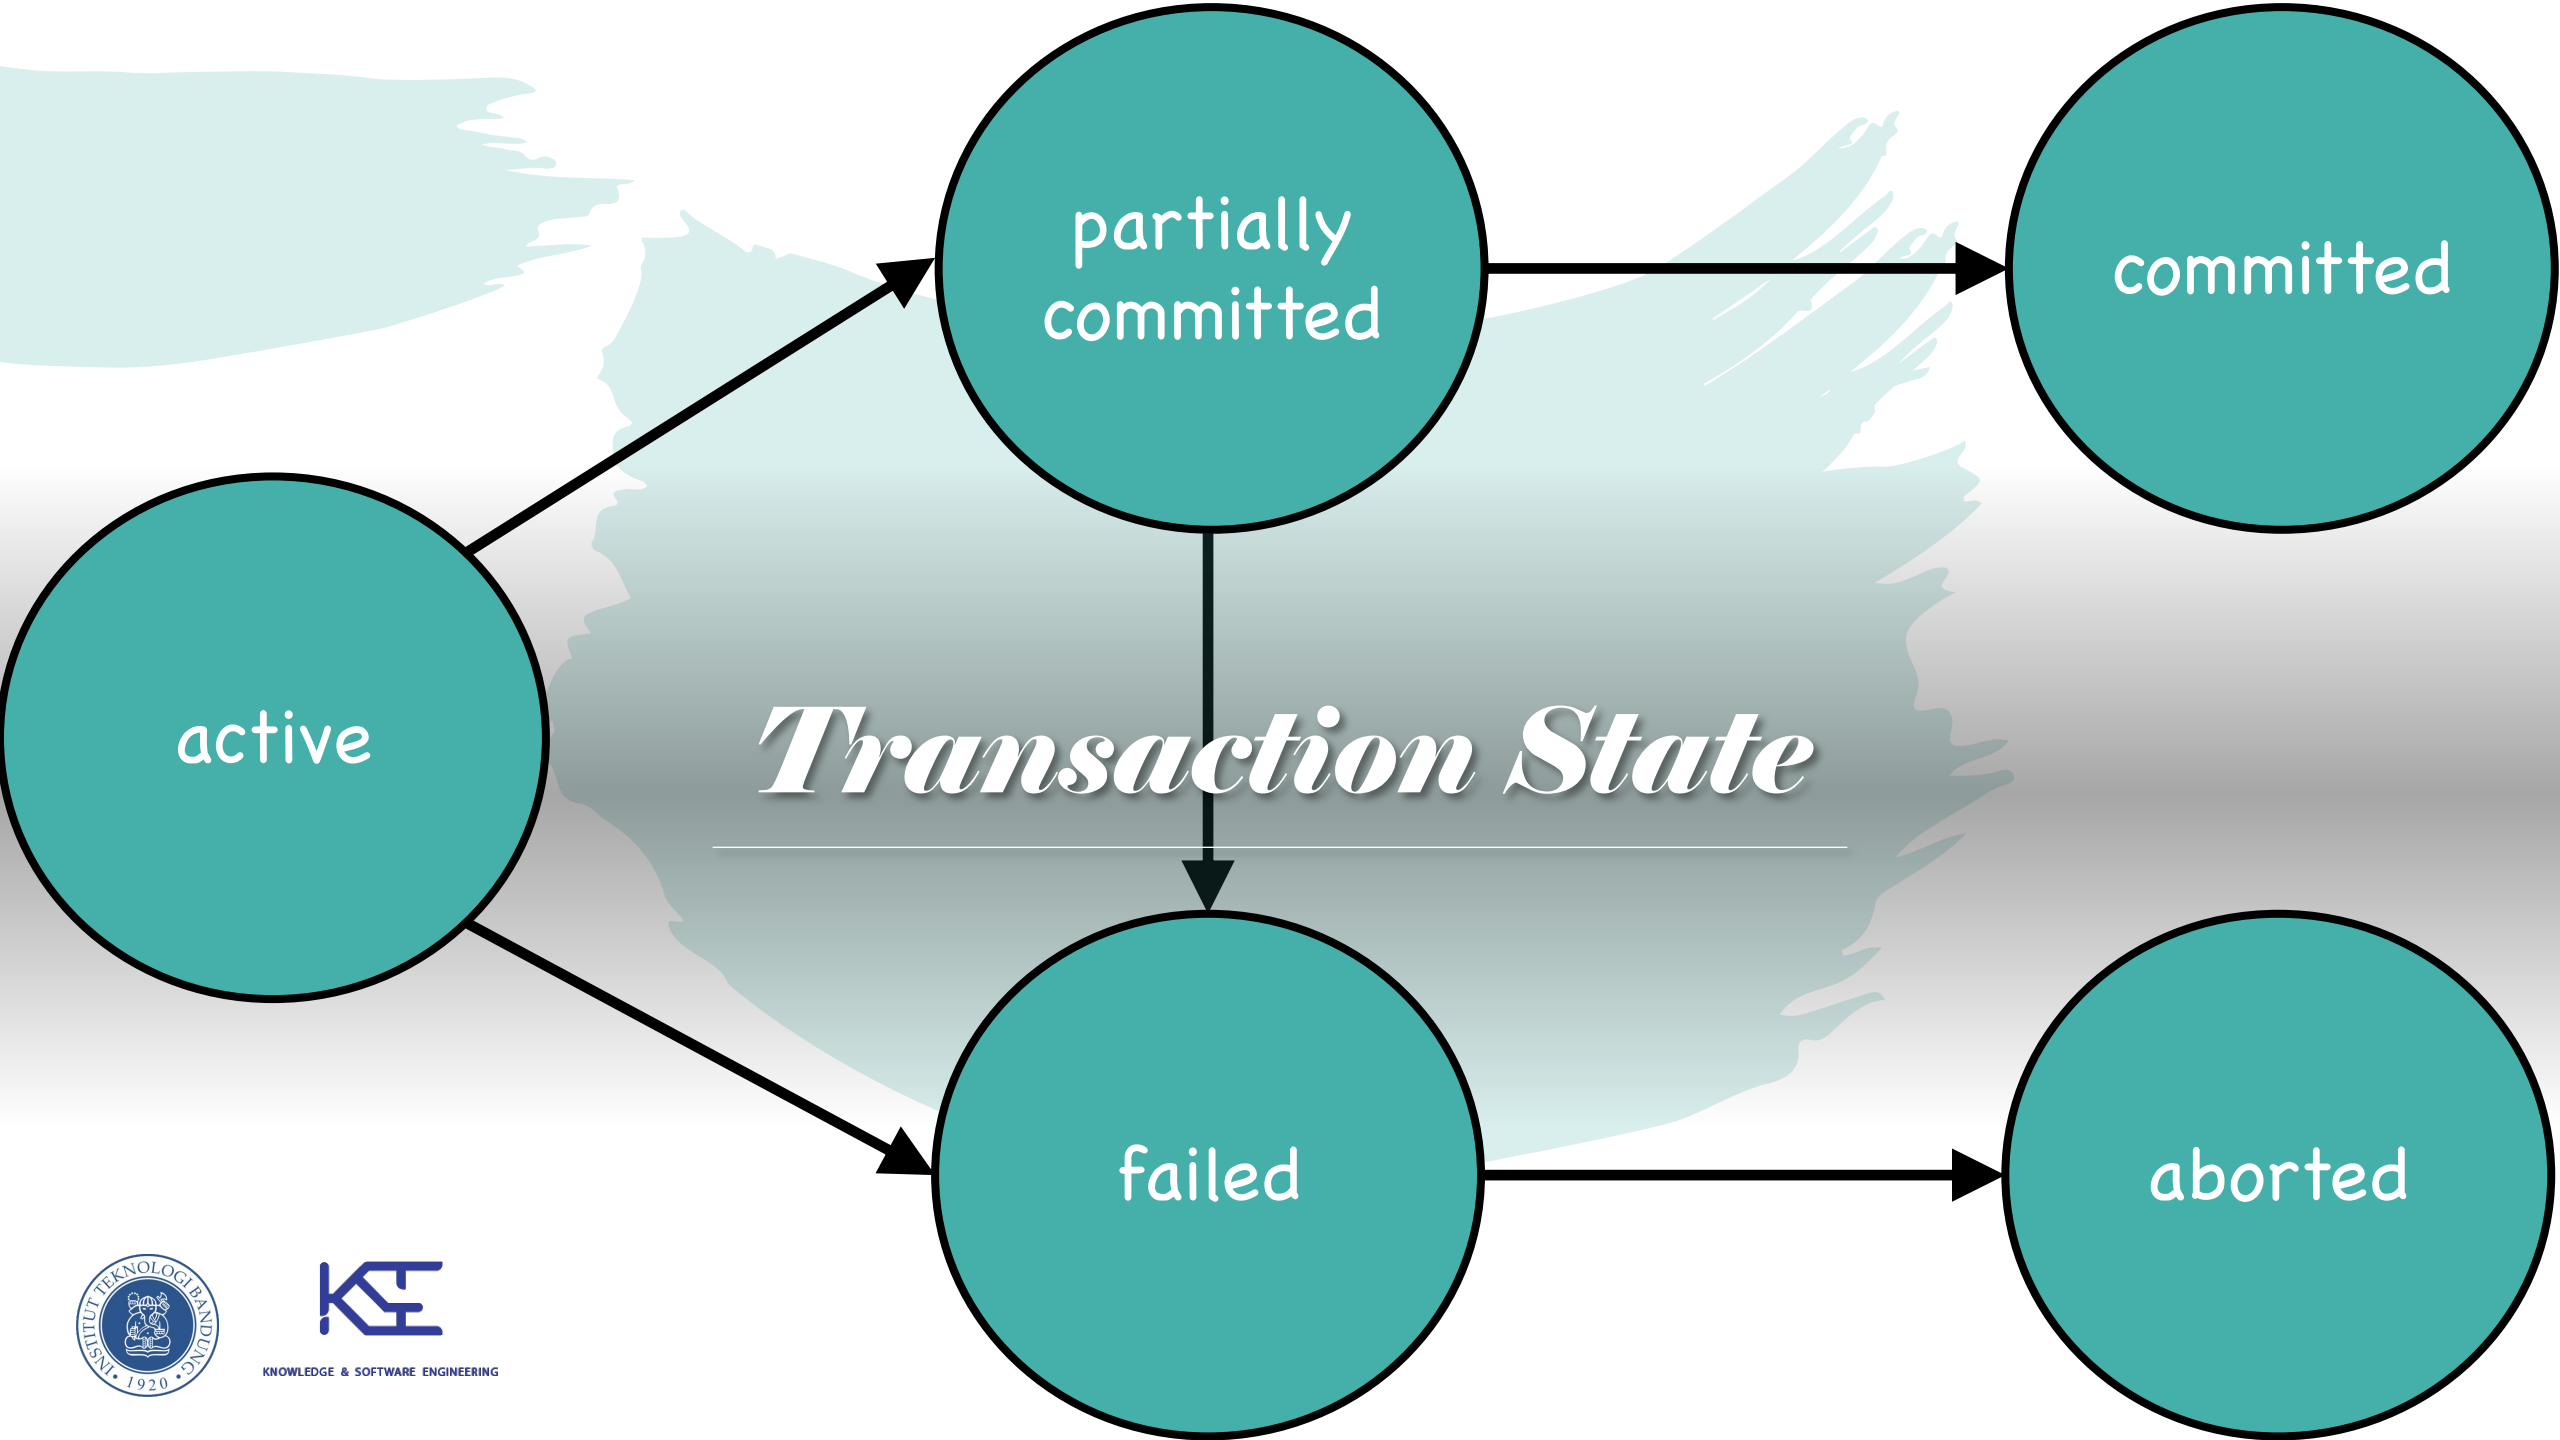
- Isolation can be ensured trivially by running transactions **serially**
 - that is, one after the other.
- However, executing multiple transactions concurrently has significant benefits, as we will see later.

ACID Properties

A **transaction** is a unit of program execution that accesses and possibly updates various data items. To preserve the integrity of data the database system must ensure:

- **Atomicity**. Either all operations of the transaction are properly reflected in the database or none are.
- **Consistency**. Execution of a transaction in isolation preserves the consistency of the database.
- **Isolation**. Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.
 - That is, for every pair of transactions T_i and T_j , it appears to T_i that either T_j finished execution before T_i started, or T_j started execution after T_i finished.
- **Durability**. After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.





Transaction State

Concurrent Executions

- Advantages are
 - **increased processor and disk utilization**, leading to better transaction throughput
 - **reduced average response time for transactions**
- **Concurrency control schemes** – mechanisms to achieve isolation

Schedules

Schedule – a sequences of instructions that specify the **chronological order** in which instructions of concurrent transactions are executed

All instructions

Preserve the order of
instructions in each
transaction

A transaction that successfully completes its execution will have a commit instructions as the last statement

A transaction that fails to successfully complete its execution will have an abort instruction as the last statement

Schedule (Serial)

Let T_1 transfer \$50 from A to B, and
 T_2 transfer 10% of the balance from A to B.



Schedule 1

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Schedule (Serial)

Let T_1 transfer \$50 from A to B, and
 T_2 transfer 10% of the balance from A to B.

Schedule 1

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Schedule 2

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Schedule 1

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Schedule 3

T_1	T_2
read (A) $A := A - 50$ write (A)	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	read (B) $B := B + temp$ write (B) commit

Schedule

*The following schedule is not a serial schedule, but it is **equivalent** to Schedule 1.*

Let T_1 transfer \$50 from A to B, and T_2 transfer 10% of the balance from A to B.

Serializability

- Basic Assumption – Each transaction preserves database consistency.
- Thus serial execution of a set of transactions preserves database consistency.
- A (possibly concurrent) schedule is **serializable** if it is equivalent to a serial schedule. Different forms of schedule equivalence give rise to the notions of:
 - **conflict serializability**
 - **view serializability**



Simplified view of transactions

- We ignore operations other than **read** and **write** instructions
- We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes.
- Our simplified schedules consist of only **read** and **write** instructions.

Conflicting Instructions

- Instructions I_i and I_j of transactions T_i and T_j respectively, **conflict** if and only if there exists some item Q accessed by both I_i and I_j , and at least one of these instructions wrote Q .
- Intuitively, a conflict between I_i and I_j forces a (logical) temporal order between them.

T1	T2
read (Q)	read (Q)

≡

T1	T2
read (Q)	read (Q)

T1	T2
read (Q)	write (Q)

≠

T1	T2
read (Q)	write (Q)

T1	T2
write (Q)	read (Q)

≠

T1	T2
write (Q)	read (Q)

T1	T2
write (Q)	write (Q)

≠

T1	T2
write (Q)	write (Q)

Conflict Serializability

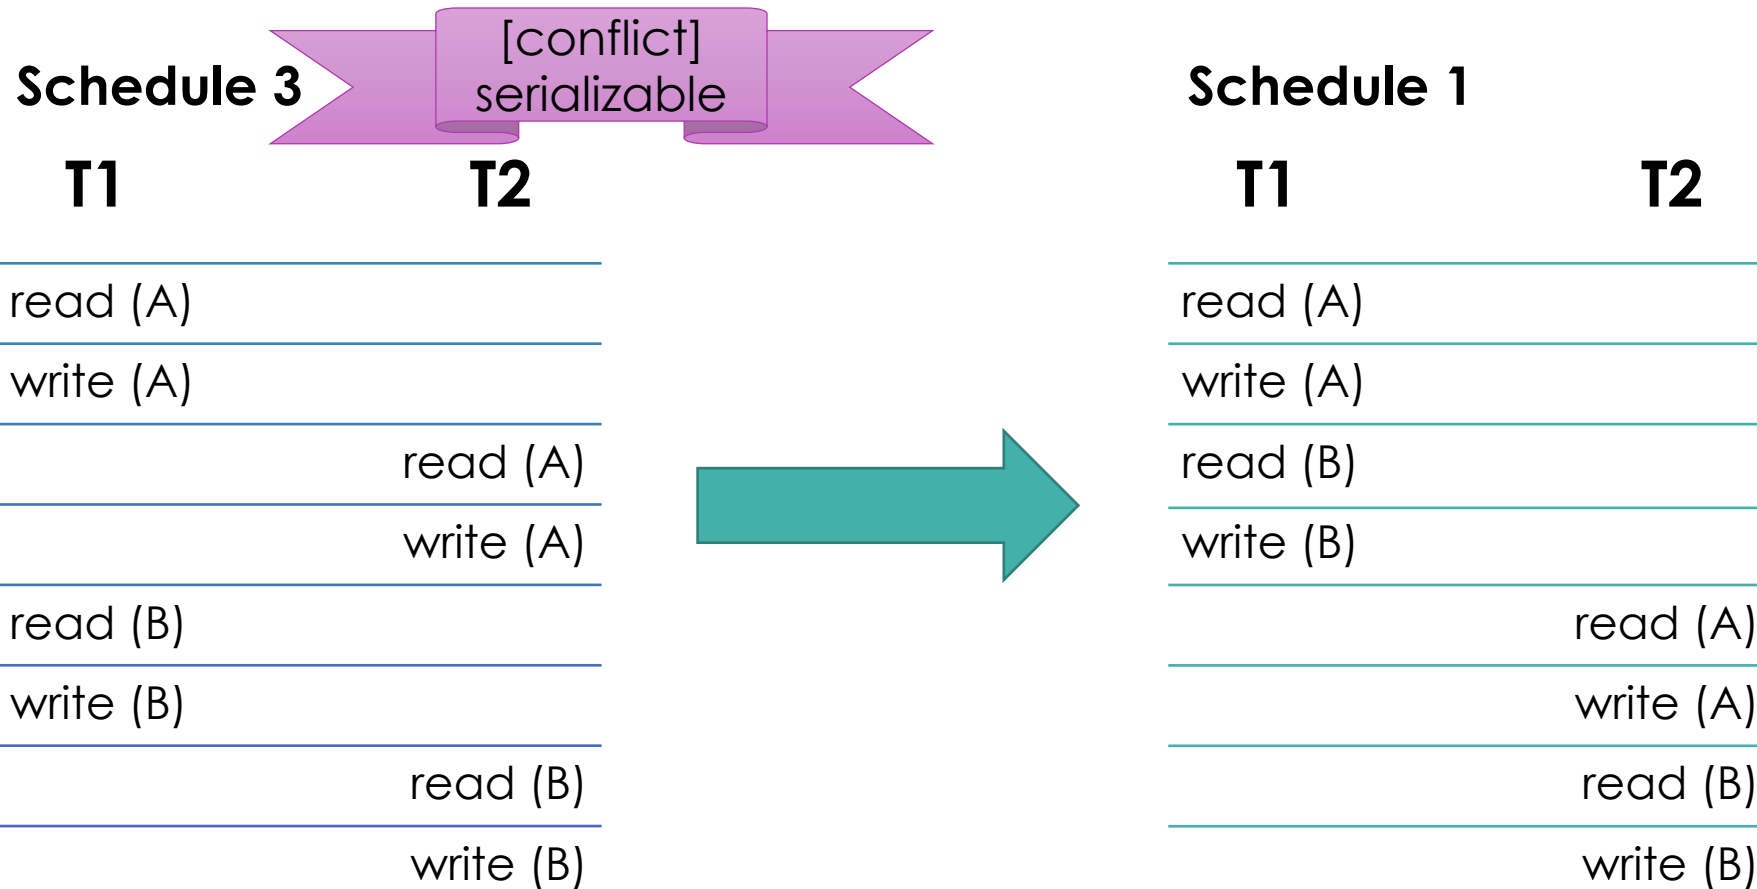
conflict equivalent schedules

schedule S can be transformed into schedule S' by a series of swaps of non-conflicting instructions

conflict serializable schedule

a schedule S that is conflict equivalent to a serial schedule

Conflict Serializability – Example 1



Conflict Serializability – Example 2

Schedule 13

T1

read (A)

write (A)

Not conflict
serializable

T2

write (A)

\neq

Schedule 11

T1

read (A)

write (A)

T2

write (A)

Schedule 12

T1

read (A)

write (A)

T2

write (A)

\neq

View Serializability

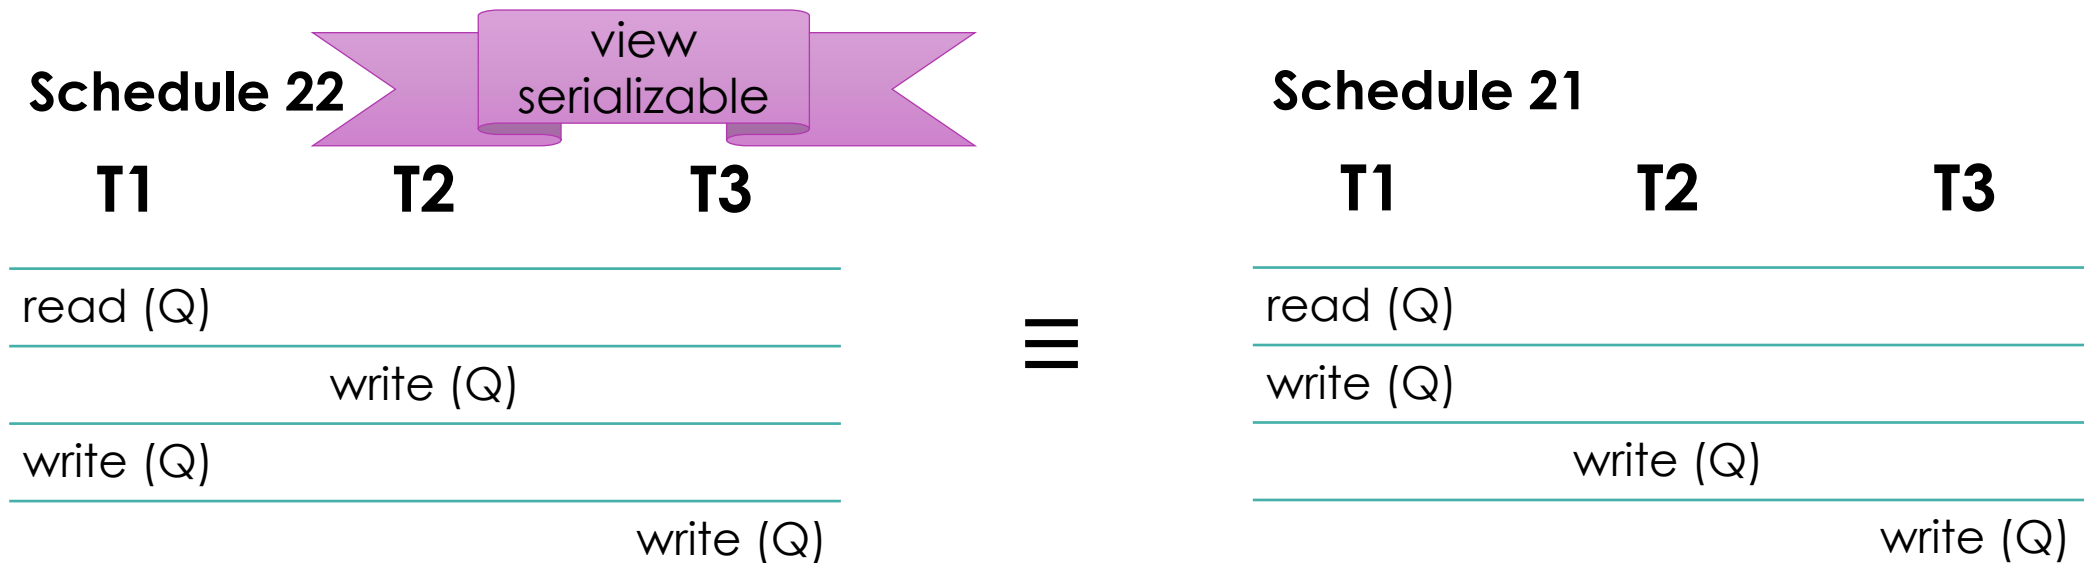
Let S and S' be two schedules with the same set of transactions.

S and S' are **view equivalent** if the following three conditions are met, for each data item Q ,

- If in schedule S transaction T_i reads the initial value of Q , then in schedule S' transaction T_i must also read the initial value of Q .
- If in schedule S transaction T_i reads the value of Q that was produced by transaction T_j , then in schedule S' transaction T_i must also read the value of Q that was produced by the same **write**(Q) operation of transaction T_j .
- The transaction (if any) that performs the final **write**(Q) operation in schedule S must also perform the final **write**(Q) operation in schedule S' .

View serializable schedule: a schedule S that is view equivalent to a serial schedule.

View Serializability – Example



T_1	T_5
read (A) $A := A - 50$ write (A)	read (B) $B := B - 10$ write (B)
read (B) $B := B + 50$ write (B)	read (A) $A := A + 10$ write (A)

Other Notions of Serializability

Determining such equivalence requires analysis of operations other than read and write

Testing for Serializability

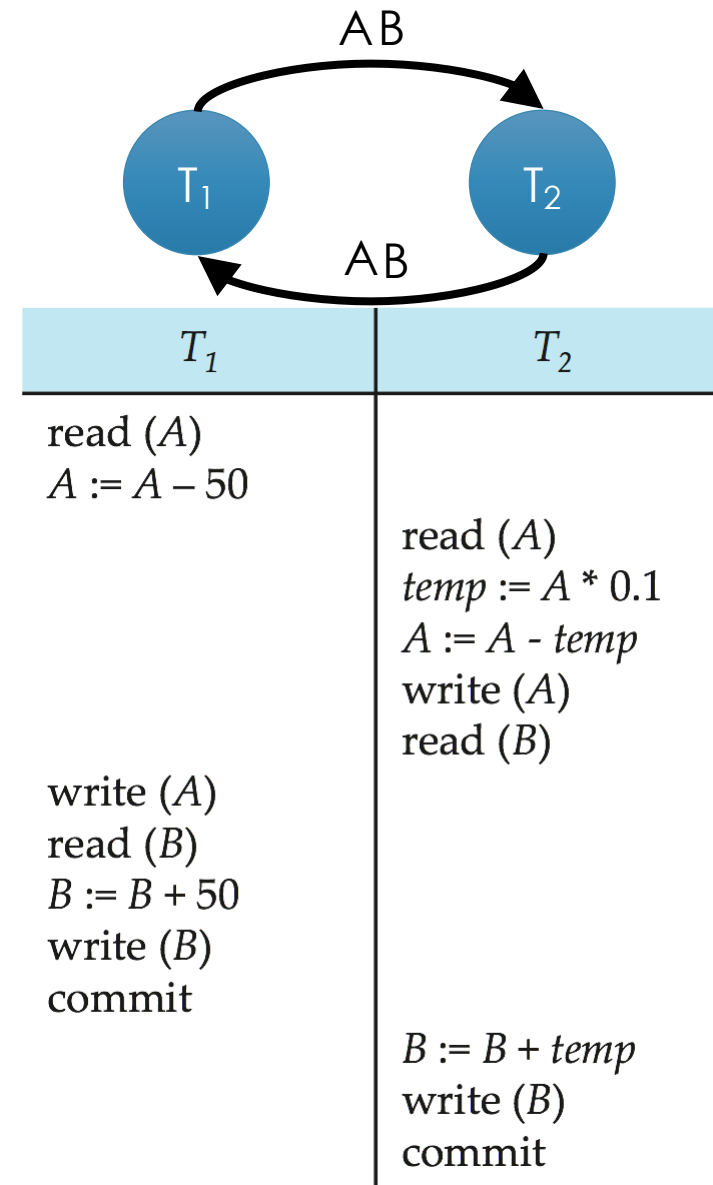
Precedence Graph

- A directed graph
- The vertices are the transactions (names).
- An arc from T_i to T_j = the two transactions conflict, and T_i accessed earlier.
- May label the arc by the item that was accessed.

Testing for Serializability

Precedence Graph




- A directed graph
- The vertices are the transactions (names).
- An arc from T_i to T_j = the two transactions conflict, and T_i accessed earlier.
- May label the arc by the item that was accessed.



Test for Conflict Serializability

A schedule is conflict serializable if and only if its **precedence graph is acyclic**.

If precedence graph is acyclic, the serializability order can be obtained by a *topological sorting* of the graph.



Test for View Serializability?

- The precedence graph test for conflict serializability cannot be used directly to test for view serializability.
 - Extension to test for view serializability has cost exponential in the size of the precedence graph.

Recoverable Schedules

If a transaction T_j reads a data item previously written by a transaction T_i , then the commit operation of T_i appears before the commit operation of T_j .

T_8	T_9
read (A)	read (A) commit
write (A)	
read (B)	
...	

Cascadeless Schedules

T_{10}	T_{11}	T_{12}
read (A) read (B) write (A)	read (A) write (A)	read (A)
abort		

- **Cascading rollback**

a single transaction failure leads to a series of transaction rollbacks.

- **Cascadeless Schedules:** no possibility of cascading rollback

- For each pair of transactions T_i and T_j such that T_j reads a data item previously written by T_i , the commit operation of T_i appears before the read operation of T_j .
- Every cascadeless schedule is also recoverable

Concurrency Control

- A database must provide a mechanism that will ensure
 - either conflict or view serializable
 - recoverable and preferably cascadeless
- Only one transaction can be executed at a time → serial schedules
→ a poor degree of concurrency
- Testing a schedule for serializability after it has executed is a little too late!
- **Goal** – to develop concurrency control protocols that will assure serializability.

Weak Levels of Consistency

TRADEOFF ACCURACY
FOR PERFORMANCE

Levels of Consistency in SQL-92

Serializable — default

Repeatable read — repeated reads must return same value.

Read committed — only committed records can be read.

Read uncommitted — even uncommitted records may be read.

Read Phenomena

- The ANSI/ISO standard SQL 92 refers to three different *read phenomena* when Transaction 1 reads data that Transaction 2 might have changed:
 - Dirty Reads
 - Non-repeatable reads
 - Phantom Reads
- Suppose we have the following data:

users		
id	name	age
1	Joe	20
2	Jill	25

Dirty Reads

- aka. Uncommitted dependency: transaction is allowed to read data from a row that has been modified by another running transaction and not yet committed.
- Level of consistency **read uncommitted** allows dirty reads

Transaction 1

```
/* Query 1 */  
SELECT age FROM users WHERE id = 1;  
/* will read 20 */
```

```
/* Query 1 */  
SELECT age FROM users WHERE id = 1;  
/* will read 21 */
```

Transaction 2

```
/* Query 2 */  
UPDATE users SET age = 21 WHERE id = 1;  
/* No commit here */
```

```
ROLLBACK; /* Lock-based DIRTY READ */
```



Non-Repeatable Reads

- During the course of a transaction, a row is retrieved twice and the values within the row differ between reads
- Level of consistency **read committed** allows non-repeatable reads

Transaction 1

```
/* Query 1 */  
SELECT * FROM users WHERE id = 1;
```

Transaction 2

```
/* Query 2 */  
UPDATE users SET age = 21 WHERE id = 1;  
COMMIT; /* in multiversion concurrency  
control, or lock-based READ COMMITTED */
```

```
/* Query 1 */  
SELECT * FROM users WHERE id = 1;  
COMMIT; /* Lock-based REPEATABLE READ */
```

Phantom Reads

- In the course of a transaction, new rows are added by another transaction to the records being read
- Level of consistency **repeatable-read** allows phantom reads

Transaction 1

```
/* Query 1 */  
SELECT * FROM users  
WHERE age BETWEEN 10 AND 30;
```

Transaction 2

```
/* Query 2 */  
INSERT INTO users(id,name,age) VALUES ( 3, 'Bob', 27 );  
COMMIT;
```

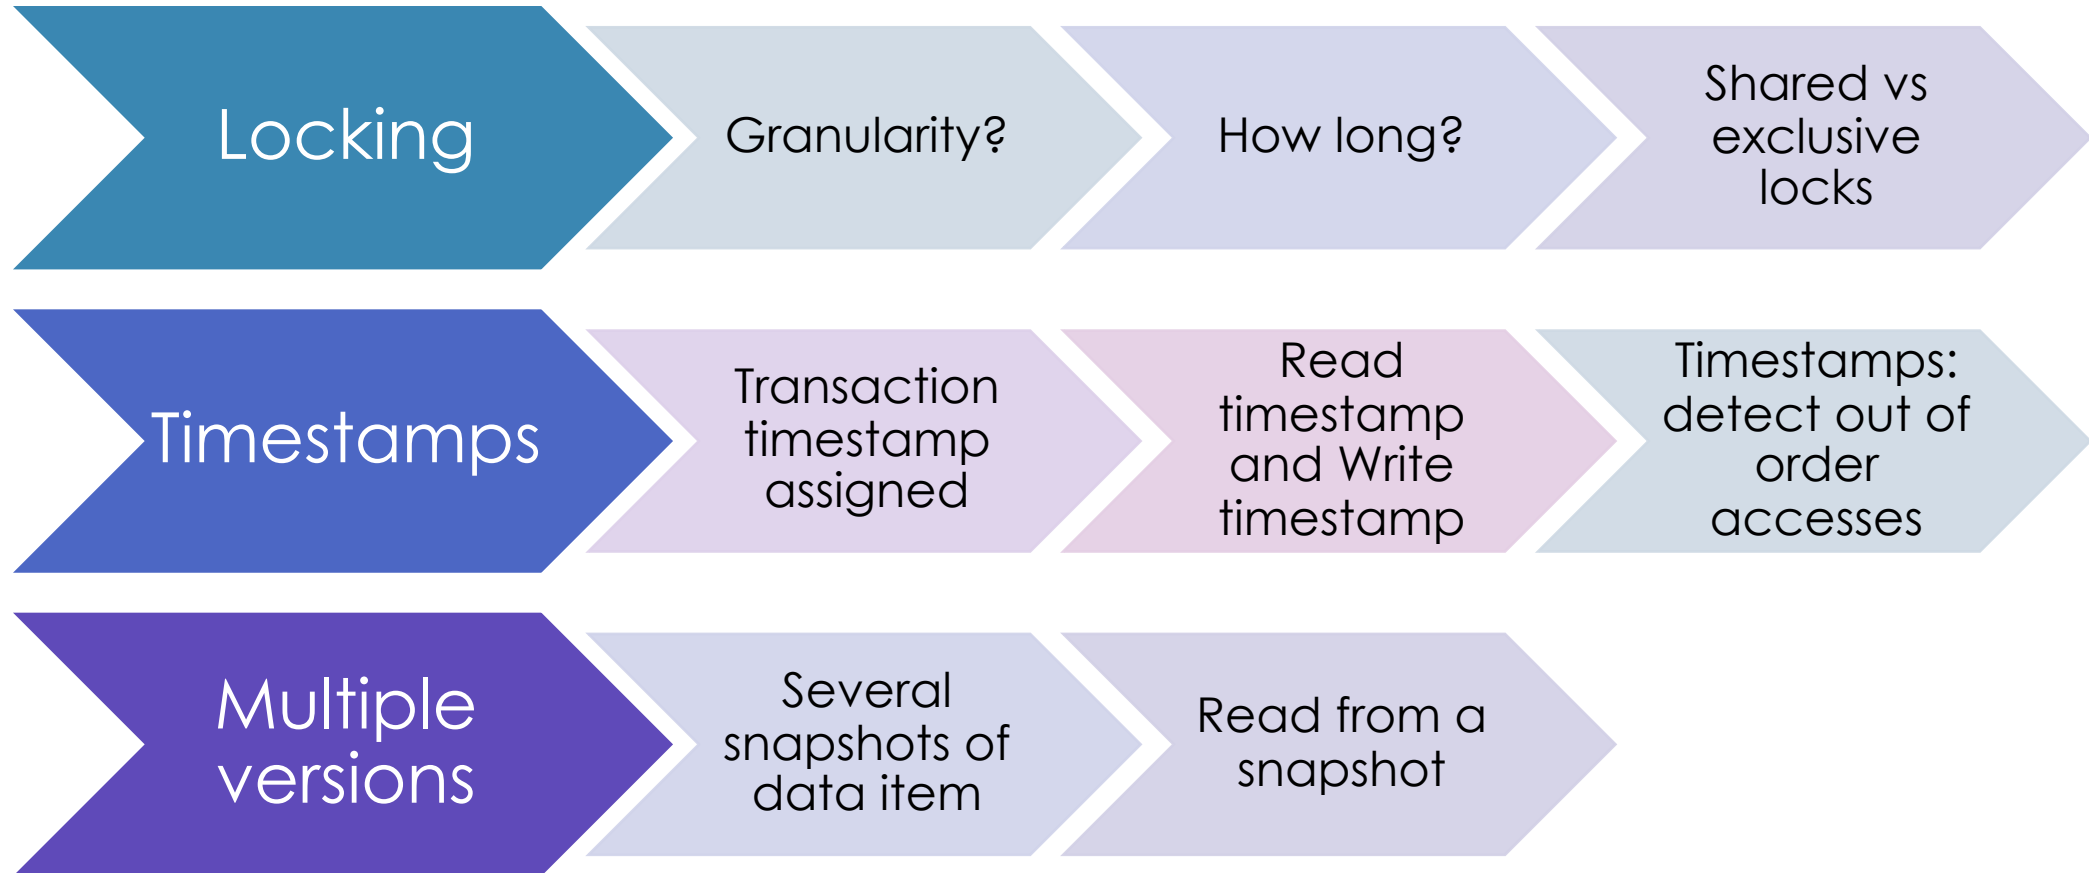
```
/* Query 1 */  
SELECT * FROM users  
WHERE age BETWEEN 10 AND 30;  
COMMIT;
```

Transaction Definition in SQL

- A transaction begins implicitly
- A transaction ends by:
 - **Commit work**
 - **Rollback work**
- Every SQL statement commits implicitly
- Isolation level can be set at database level
- Isolation level can be changed at start of transaction



Implementation of Isolation Levels



Transactions as SQL Statements

Transaction 1

```
select ID, name from instructor  
where salary > 90000
```

Transaction 2

```
insert into instructor  
values ('11111', 'James', 'Marketing', 100000)
```

Transaction 3
(with Wu's salary
= 90000)

```
update instructor set salary = salary * 1.1  
where name = 'Wu'
```

- Key idea: Detect “**predicate**” conflicts → “**predicate locking**”



End of Chapter



KNOWLEDGE & SOFTWARE ENGINEERING