


IF3140 – Sistem Basis Data

Concurrency Control:

- Introduction
- Lock-based Protocol



KNOWLEDGE & SOFTWARE ENGINEERING



Summer

Silberschatz, Korth, Sudarshan:
“Database System Concepts”,
7th Edition

- Chapter 18:
Concurrency Control

Objectives

Students are able to:

- Explain the effect of different isolation levels on the concurrency control mechanisms
- Choose the proper isolation level for implementing a specified transaction protocol

Outline

- **Lock-Based Protocols**
 - 2-phase locking
 - Graph-based protocols
 - Deadlock handling
 - Multiple granularity
- **Timestamp-Based Protocols**
- **Validation-Based Protocols**
- **Insert and delete operations**
- **Multiversion Schemes**
 - MV Timestamp ordering
 - MV 2-phase locking
 - Snapshot isolation
- **Weak levels of consistency**



Lock-Based Protocols

- A lock is a mechanism to control concurrent access to a data item
- Data items can be locked in two modes :
 1. **exclusive** (*X*) mode. Data item can be both read as well as written. X-lock is requested using **lock-X** instruction.
 2. **shared** (*S*) mode. Data item can only be read. S-lock is requested using **lock-S** instruction.
- Lock requests are made to concurrency-control manager. Transaction can proceed only after request is granted.



Lock-Based Protocols (Cont.)

- Lock-compatibility matrix**

	S	X
S	true	false
X	false	false

- A transaction may be granted a lock on an item if the requested lock is compatible with locks already held on the item by other transactions
- Any number of transactions can hold shared locks on an item,
- But if any transaction holds an exclusive on the item no other transaction may hold any lock on the item.

Schedule With Lock Grants


- Grants omitted in rest of chapter
 - Assume grant happens just before the next instruction following lock request
- This schedule is not serializable (why?)
- A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks.
- Locking protocols enforce serializability by restricting the set of possible schedules.

T_1	T_2	concurrency-control manager
lock-X(B)		grant-X(B, T_1)
read(B)		
$B := B - 50$		
write(B)		
unlock(B)		
	lock-S(A)	
	read(A)	grant-S(A, T_2)
	unlock(A)	
	lock-S(B)	
		grant-S(B, T_2)
	read(B)	
	unlock(B)	
	display($A + B$)	
lock-X(A)		grant-X(A, T_1)
read(A)		
$A := A + 50$		
write(A)		
unlock(A)		



The Two-Phase Locking Protocol

- A protocol which ensures conflict-serializable schedules.
- Phase 1: **Growing Phase**
 - Transaction may obtain locks
 - Transaction may not release locks
- Phase 2: **Shrinking Phase**
 - Transaction may release locks
 - Transaction may not obtain locks
- The protocol assures serializability. It can be proved that the transactions can be serialized in the order of their **lock points** (i.e., the point where a transaction acquired its final lock).



Deadlock

- Consider the partial schedule

T_3	T_4
lock-X(B) read(B) $B := B - 50$ write(B)	
	lock-S(A) read(A) lock-S(B)
lock-X(A)	

- Neither T_3 nor T_4 can make progress — executing **lock-S(B)** causes T_4 to wait for T_3 to release its lock on B , while executing **lock-X(A)** causes T_3 to wait for T_4 to release its lock on A .
- Such a situation is called a **deadlock**.
 - To handle a deadlock one of T_3 or T_4 must be rolled back and its locks released.

Deadlock (Cont.)

- The potential for deadlock exists in most locking protocols. Deadlocks are a necessary evil.
- **Starvation** is also possible if concurrency control manager is badly designed. For example:
 - A transaction may be waiting for an X-lock on an item, while a sequence of other transactions request and are granted an S-lock on the same item.
 - The same transaction is repeatedly rolled back due to deadlocks.
- Concurrency control manager can be designed to prevent starvation.



The Two-Phase Locking Protocol (Cont.)

- Two-phase locking *does not* ensure freedom from deadlocks
- Extensions to basic two-phase locking needed to ensure recoverability or freedom from cascading roll-back
 - **Strict two-phase locking:** a transaction must hold all its exclusive locks till it commits/aborts.
 - Ensures recoverability and avoids cascading roll-backs
 - **Rigorous two-phase locking:** a transaction must hold *all* locks till commit/abort.
 - Transactions can be serialized in the order in which they commit.
- Most databases implement rigorous two-phase locking, *but refer to it as simply two-phase locking*



The Two-Phase Locking Protocol (Cont.)

- Two-phase locking is not a necessary condition for serializability
 - There are conflict serializable schedules that cannot be obtained if the two-phase locking protocol is used.
- In the absence of extra information (e.g., ordering of access to data), two-phase locking is necessary for conflict serializability *in the following sense*:
 - Given a transaction T_i that does not follow two-phase locking, we can find a transaction T_j that uses two-phase locking, and a schedule for T_i and T_j that is not conflict serializable.

T_1	T_2
lock-X(B)	
read(B)	
$B := B - 50$	
write(B)	
unlock(B)	lock-S(A)
	read(A)
	unlock(A)
	lock-S(B)
	read(B)
	unlock(B)
	display($A + B$)
lock-X(A)	
read(A)	
$A := A + 50$	
write(A)	
unlock(A)	

©Silberschatz et.al. (2019)

[modified by Tim Pengajar IF3140 Semester 1 2024/2025]



Locking Protocols

- Given a locking protocol (such as 2PL)
 - A schedule S is **legal** under a locking protocol if it can be generated by a set of transactions that follow the protocol
 - A protocol **ensures** serializability if all legal schedules under that protocol are serializable



Exercise 9.1

Contoh alternatif yg akan deadlock

T1	T2
L-S(A)	L-S(B)
L-X(B)	L-X(A) → wait T1
R(A)	
R(B)	
	R(B)
	R(A)

Contoh eksekusi yg ga akan deadlock, TP jd serial karena T2 harus menunggu T1 all bekes dulu

T1	T2
L-S(A)	
L-X(B)	
R(A)	
R(B)	
	L-S(B)
	L-X(A)
	R(B)
	R(A)
	W(B)
	UL(B)
	UL(A)
	W(A)
	UL(A)
	UL(B)

↳ harus wait soalnya di lock di T1.

Consider the following two transactions:

T1: read (A)

read (B)

$B := B + 0.1 * A$

write (B)

T2: read (B)

read (A)

$A := A - 0.05 * B$

write (A)

- Add lock and unlock instructions to both transactions so that they follow the two-phase locking protocol.
- Can the execution of these two transactions result in a deadlock?



KNOWLEDGE & SOFTWARE ENGINEERING

Solution

©Silberschatz et.al. (2019)

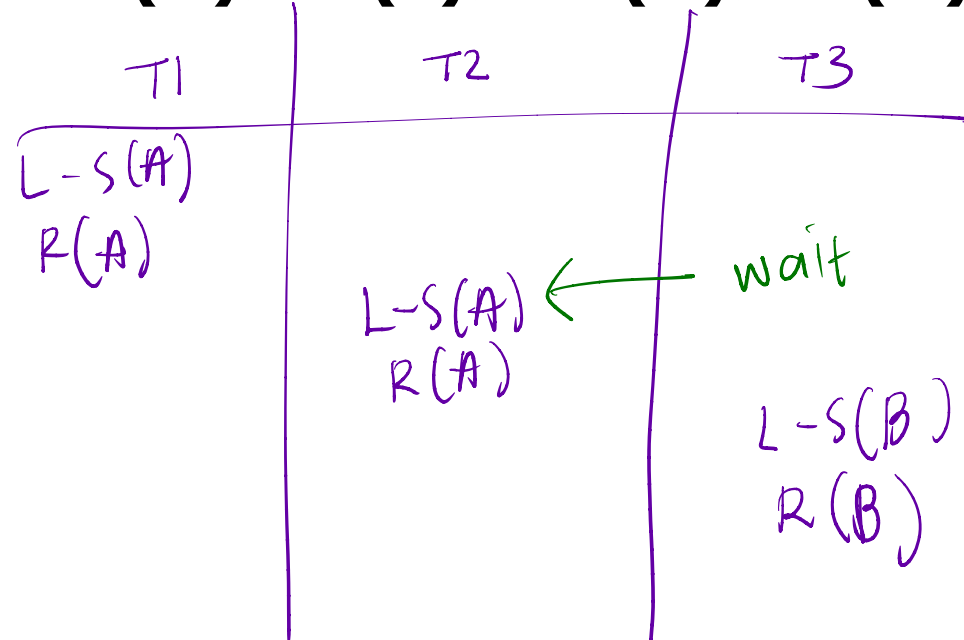
[modified by Tim Pengajar IF3140 Semester 1 2024/2025]

Exercise 9.2

NOTE: automatic bisa ubah dr $L(S)$ ke $L(X)$ → read
→ read & write

Is the following schedule a two-phase locking (2PL) schedule (legal under 2PL protocol)?

R1(A); R2(A); R3(B); W1(A); R2(C); R2(B); W2(B); W1(C);



Lock Conversions

- Two-phase locking protocol with lock conversions:
 - Growing Phase:
 - can acquire a lock-S on item
 - can acquire a lock-X on item
 - can **convert** a lock-S to a lock-X (**upgrade**)
 - Shrinking Phase:
 - can release a lock-S
 - can release a lock-X
 - can convert a lock-X to a lock-S (**downgrade**)
- This protocol ensures serializability



Automatic Acquisition of Locks

- A transaction T_i issues the standard read/write instruction, without explicit locking calls.
- The operation **read**(D) is processed as:
 - if** T_i has a lock on D
 - then**
 - read(D)
 - else begin**
 - if necessary wait until no other transaction has a **lock-X** on D
 - grant T_i a **lock-S** on D ;
 - read(D)
 - end**

Automatic Acquisition of Locks (Cont.)

- The operation **write**(D) is processed as:
 - if** T_i has a **lock-X** on D
 - then**
 - write(D)
 - else begin**
 - if necessary wait until no other trans. has any lock on D ,
 - if T_i has a **lock-S** on D
 - then**
 - upgrade** lock on D to **lock-X**
 - else**
 - grant T_i a **lock-X** on D
 - write(D)
 - end;**
- **All locks are released after commit or abort**



Exercise 9.3

Instructions from T1, T2, and T3 arrive in the following order.

**R1(A); R2(A); R3(B); W1(A); R2(C); R2(B); C3; W2(B); C2;
W1(C); C1;**

What is the final schedule if the 2-phase locking with automatic acquisition of locks is implemented by CC Manager?

